

PM/Product Learning Digest Agent

Backend AI Engineering Capstone · Sibgha Mursaleen

AGENT OVERVIEW

Checks a set of trusted project management and product management sources for new articles, classifies each one by skill area (e.g. Prioritization, Agile, Stakeholder Management, Roadmapping), and saves the sorted links — title, source, and original URL — into a running reading list. No summarizing or rewriting of source content; the agent's only judgment call is the skill-area classification, and every entry links back to the original article.

User: Solo use (me).

Frequency: Run on demand or on a schedule, a few times a week.

DEVIATION FROM FL-06 SPEC

Original spec: a Job Application Tracker & Follow-Up Agent (Gmail + Google Sheets, OAuth-based).

Change: pivoted to a PM/Product Learning Digest Agent.

Reason: self-hosted n8n made Google OAuth setup a heavy, low-value time cost relative to a 10-hour build. This job needed only RSS feeds and a local file — no OAuth at all — while staying equally useful and better matched to a current, active goal (building PM/product skills).

TOOLS & DATA — ACCESS PLAN

Tool	Purpose	Access Plan
RSS feeds	Source new articles from trusted PM/product blogs	n8n RSS Feed Read node — no authentication required
LLM (OpenRouter)	Classify each article into a skill-area tag	Existing API key from FL-04
Local file	Store the deduplicated, tagged reading list (reading-list.md) and a seen-links log (seen-links.json)	n8n Read/Write File node — local filesystem, no authentication required

DRAFT INSTRUCTIONS

- On trigger: fetch latest items from each configured RSS feed.
- Compare fetched items against seen-links.json; discard anything already logged.
- For each new item, classify it into one skill-area tag using the LLM (single word/phrase, not a summary).
- Append new items to reading-list.md, grouped under their skill-area heading, with title, source, link, and date.
- Update seen-links.json with the newly processed links.
- If a feed is unreachable, skip it and continue processing the remaining feeds — one broken source must not stop the run.

EVAL CASES

- New articles exist → correctly appended with title, source, link, date, and skill-area tag.

2. No new articles → nothing appended, file unchanged.
3. Same article appears again on a later run → correctly skipped, not duplicated.
4. One feed is broken or unreachable → remaining feeds still process; run completes.
5. Two feeds cover the same article → handled without duplicate or malformed entries.
6. An article's topic is ambiguous → agent still assigns its best single tag rather than failing the run.

RISKS & GUARDRAILS

Must never: summarize, rewrite, or reproduce article content — only title, source, link, date, and a one-word/phrase skill tag are stored.

Must always: keep the original source link visible and unaltered, so every entry can be independently verified.

Failure mode to design for: a broken or slow feed must not block the rest of the run — skip and continue.

BUILD LOG

Filled in as the build progresses — real iteration, not a retroactive clean story.

Date	What I tried	What broke	What I changed / cut, and why
	Set up Google OAuth for Gmail + Sheets per FL-06 spec	Self-hosted n8n needed a manually configured Google Cloud OAuth client; high setup cost for a 10-hour build	Cut Gmail/Sheets entirely; pivoted to an RSS + local-file agent needing no OAuth (see Deviation section above)

PLATFORM

n8n (self-hosted, local), continuing from FL-06. RSS Feed Read, HTTP Request (LLM classification), and Read/Write File nodes cover the full loop without any OAuth-based credentials.

SUBMISSION CHECKLIST

- Working agent — completes the full loop end to end without mid-run hand-editing.
- Build log — filled in above with real entries as the build happens.
- Raw run capture — ~2 minute unedited screen recording of one successful end-to-end run, from trigger to result.